

Learning Transferable Graph Exploration

Authors: Hanjun Dai et al.

Presented by Qinghai Zhou

Outline

- Background and Problem Formulation
- Methodology
- Evaluation

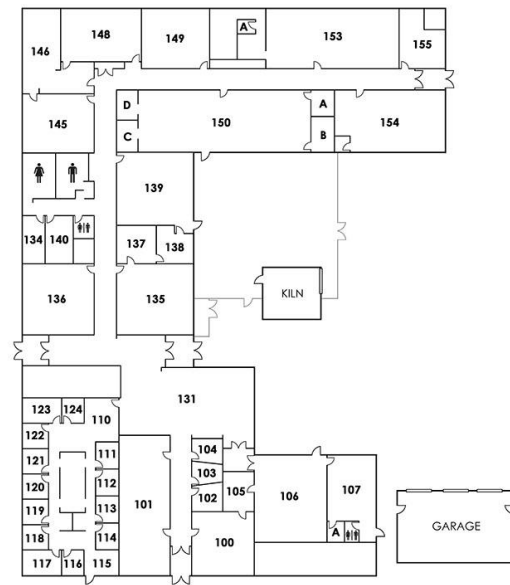
Background

State-space Coverage Problem

Goal: given an environment, efficiently reach as distinct/unique states as possible

Applications:

- model checking: design test inputs to expose as many potential errors as possible
- active map building: construct a map of unknown environment efficiently
- exploration in reinforcement learning in general



Common Approaches: Undirected Exploration

Key idea: randomly choose states to visit / actions to take

Examples:

1. Random walk based method on graph
 - Cover time depends on graph structure ($O(n \log n)$, $O(n^3)$)
2. Epsilon-greedy selection
 - selecting random action with probability epsilon
 - preventing being locked on sub-optimal action

Directed Exploration

Key idea: optimize a certain objective that encourages exploration/coverage

Examples:

1. Upper Confidence Bound (UCB) for Bandit Problems:
 - in addition to maximizing the reward, encourage exploring unselected actions
2. Reinforcement Learning (RL)
 - pseudo-count (similar to UCB): rewards change in state density estimates
 - information gain: take actions from which you learn about the environment (reduces entropy)
 - predictive error: encourage actions that lead to unpredictable outcome (for instance, unseen states)

Graph Exploration

Goal: an efficient exploration strategy to reach as many vertices as possible

- effectiveness of random walk greatly depends on the graph structure

Question to be answered:

Given the distribution of graphs in training time, can the algorithm learn an efficient covering/exploration strategy?

Formulation

Several key components in RL:

- Environment: for agent to interact with
- Action: state to cover/explore
- Reward: feedback from the environment

Formulation

Environment: Graph-structured state-space


- at time t , the agent observes a graph $G_{t-1} = \{E_{t-1}, V_{t-1}\}$, and a coverage mask $\epsilon_{t-1} : V_{t-1} \rightarrow \{0, 1\}$ indicating the nodes' exploration status so far
- the agent takes an action a_t and receives a new graph G_t
- Budget: number of steps/actions allowed to take Same dist. as training

Goal:

learn exploration strategy such that given an unseen environment, the agent can efficiently visit as many unique nodes as possible

Reward (maximize the number of visited nodes)

Cumulative rewards: $\max_{\{a_1, a_2 \dots a_t\}} \sum_{v \in V_t} \frac{c_t(v)}{|V|}$

 Coverage mask

Per-step reward: $r_t = \sum_{v \in V_t} \frac{c_t(v)}{|V|} - \sum_{v \in V_{t-1}} \frac{c_{t-1}(v)}{|V|}$

They are equivalent $\sum_{v \in V_0} c_0(v) = 0$

Objective

$$\max_{\{\theta_1, \theta_2 \dots \theta_t\}} \mathbb{E}_{\mathcal{G} \sim \mathcal{D}} \left[\sum_{t=1}^T \mathbb{E}_{a_t^{\mathcal{G}} \sim \pi(a|h_t^{\mathcal{G}}, \theta_t)} [r_t^{\mathcal{G}}] \right]$$

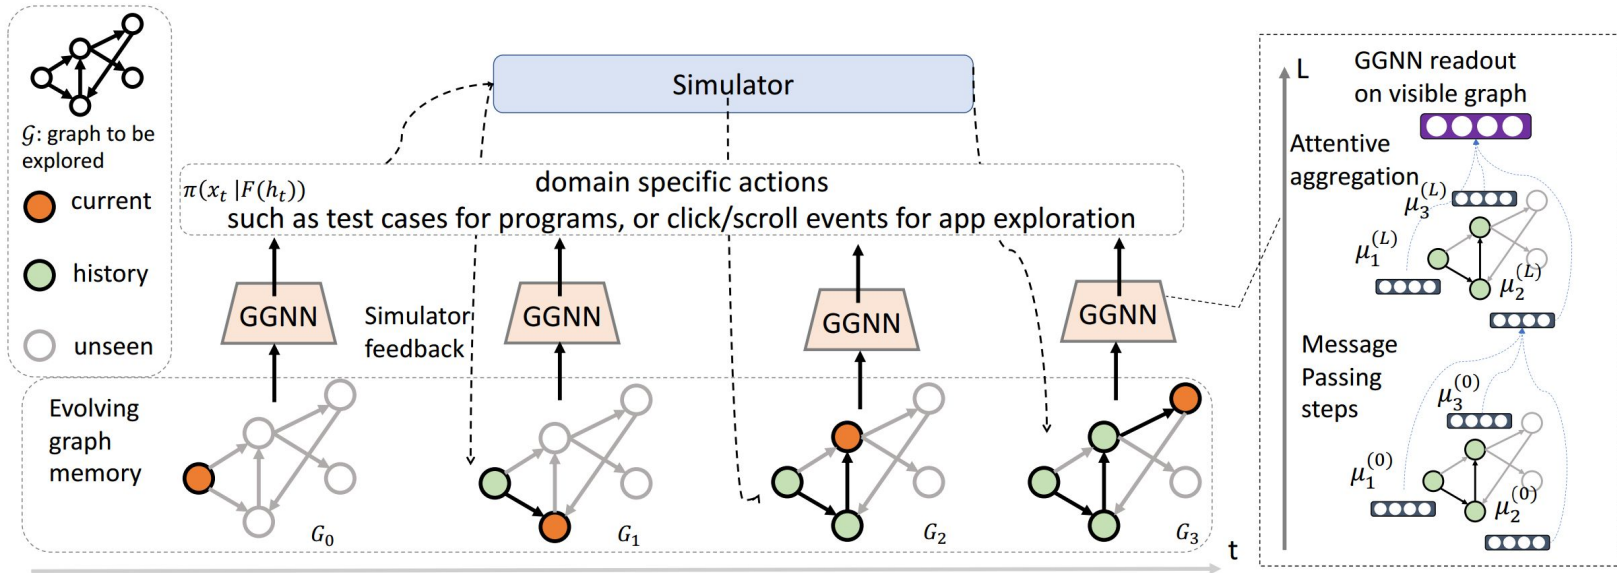
- $h_t = \{(a_i, G_i, c_i)\} (i=1, \dots, t)$ is the exploration history
- $\pi(a|h_t, \theta_t)$ is the action policy at time t
- \mathcal{D} is the environment distribution

Representing Exploration History

Graph representation:

- use graph neural networks to learn a representation (one-bit information c_t)
- starting from node $\mu_v^{(0)}$, update representation via message passing: $\mu_v^{(l+1)} = f(\mu_v^{(l)}, \{e_{uv}; \mu_u^{(l)}\})$, ($f(\cdot)$: GGNN model)
- apply attention weighted-aggregation
- learned via unsupervised link prediction

Representing Exploration History




History representation: summarize representation up to the current step

Auto-regressive $F(h_t) = \text{LSTM}(F(h_{t-1}, g(G_t, c_t))) \rightarrow \text{encoder}$

Evaluation


Two settings: (1) unknown graph; (2) known but complex (program testing)



Erdos-Renyi Random Graph:

blue node indicates starting point; darker colors represent more visit counts

Evaluation



Full Maze Random RandDFS GMETAEXP


2D Maze

- given fixed budget ($T = 36$), the agent is trained to traverse the 6x6 maze as much as possible
- test on held-out mazes from the same distribution

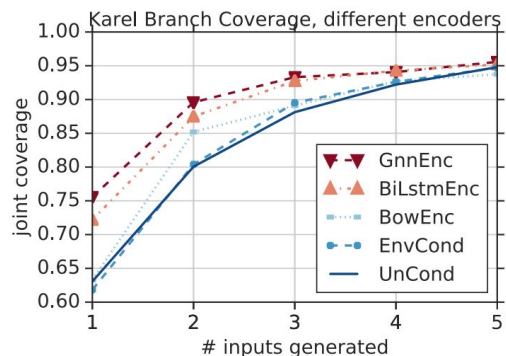

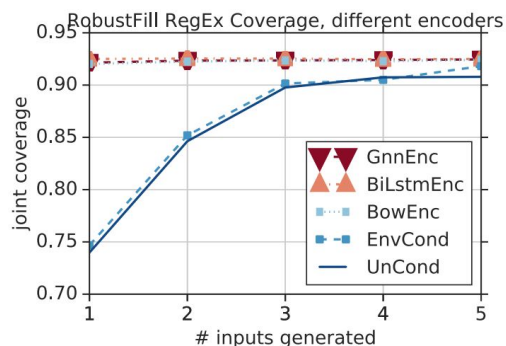
Evaluation

n

(a) Program coverage results



(b) History encoding



**Program testing
(coverage guided
fuzzing):**

Propose test cases to
cover as many code
branches as possible

Limitation:

- requires reasonable large amount of training data
- cannot scale to large programs

Possible extension:

- RL-based dynamic graph representation

Please check the paper for more details about method and experiment.

**Thank
you!
Q & A**