
Disentangling Degree-related Biases and Interest for Out-of-distribution Generalized Directed Network Embedding

2023.02.02


Hyunsik Yoo

* This is a joint work with Prof. Sang-Wook Kim and Dr. Yeon-Chang Lee at **Hanyang Univ.**, and Prof. Kijung Shin at **KAIST**

Background: Network Embedding (NE)

□ Represents nodes in a given network as **low-dimensional vectors that preserve the structural properties of the network**

■ *e.g.*, proximity between nodes



□ The learned embeddings can be used as informative features of nodes in **various downstream network mining tasks**

■ Link prediction → Our focus


■ Node clustering/classification

■ Recommendation

Background: Directed Network

□ A directed network

- A directed edge from node i to j expresses an **asymmetric relationship** (or proximities) between two nodes
- A toy example on Instagram




□ To capture such asymmetric relationships accurately, various **directed network embedding (DNE) methods** have been proposed

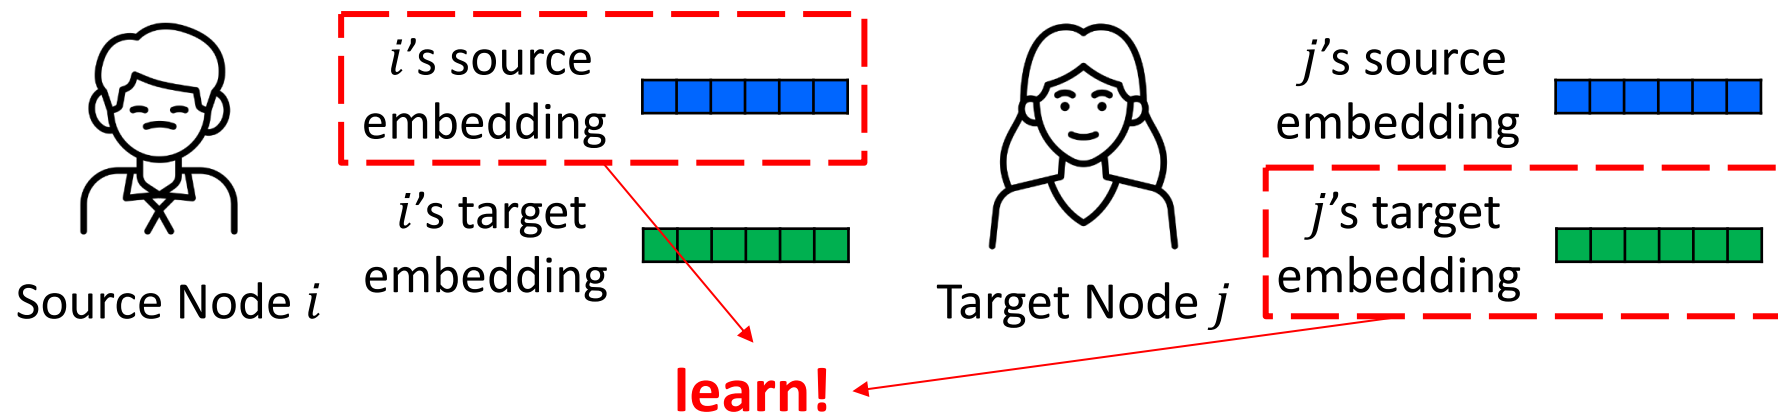
- APP [AAAI'17]
- ATP [AAAI'19]
- NERD [ECML-PKDD'19]
- GVAE [CIKM'19]
- DiGCN [NeurIPS'20]
- MagNet [NeurIPS'21]
- DGGAN [AAAI'21]

Background: Directed Network (cont'd)

□ Given a directed edge from i to j ,



- Distinguish the **source node i** and the **target node j** according to their roles
- Learn a **source embedding** and a **target embedding**, which preserve the node's properties as sources and targets



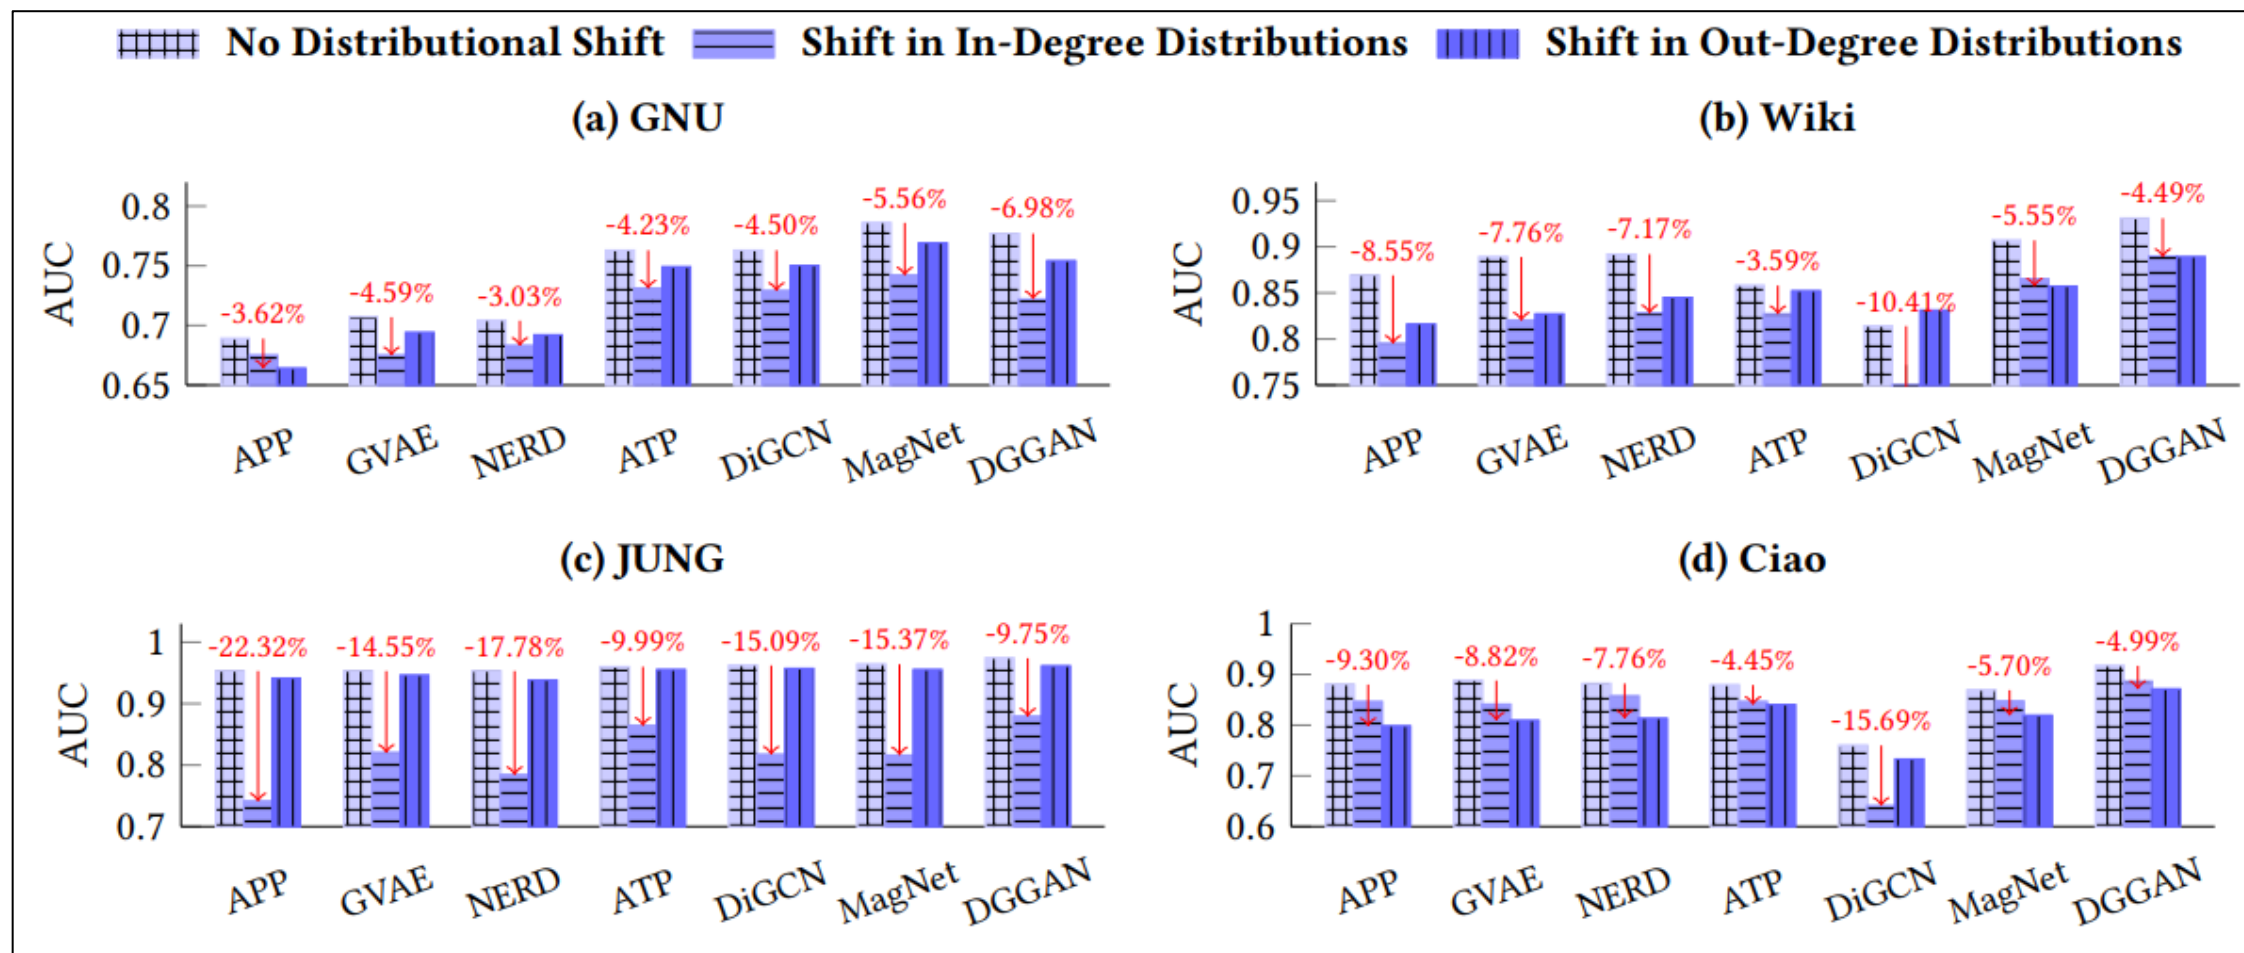
Motivation

- ❑ Existing DNE methods lack out-of-distribution (OOD) generalization abilities against degree-related distributional shifts
 - They assume that, in link prediction, the degree distribution of the training and test data are identical
- ❑ However, in real-world scenarios, degree-related distributional shifts occur frequently → ruining the identical distribution (ID) assumption!
 - Preferential attachment
 - Fitness model: it is also common that dominant hubs are overtaken by “new kids on the block” with higher fitness
 - ❑ e.g., Google passed established search engines, such as Alta Vista

Motivation (cont'd)

□ Link prediction accuracy of the existing methods in ID / non-ID settings

- The accuracies of all methods significantly degrade in the non-ID settings compared to the ID settings




Proposed Method

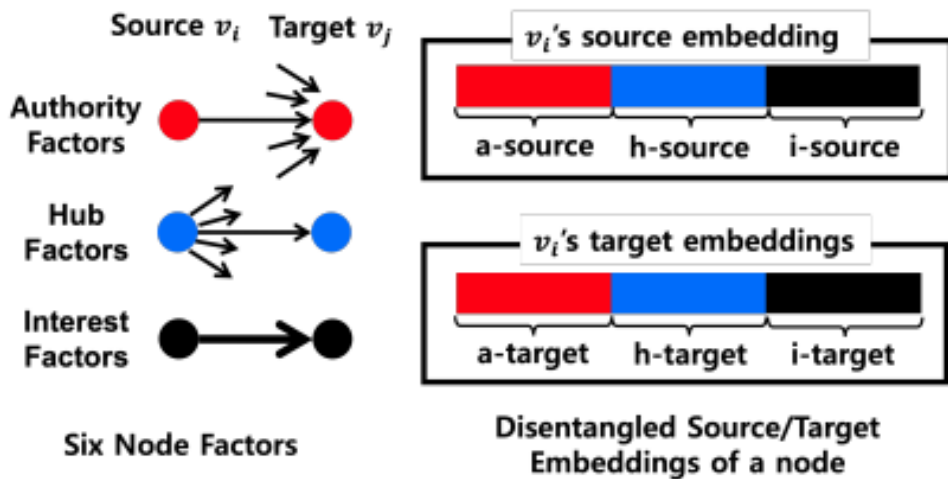
- Our idea: model and exploit biases related to node degrees for robustness against degree-related distributional shifts in DNE

- Propose **ODIN** (Out-of-Distribution Generalized Directed Network Embedding), which is designed to answer the following questions:
 1. *How to model the formation of each directed edge?*
 - Define six node factors that can influence the formation of a directed edge from source to target
 2. *How to leverage such modeled factors for learning OOD generalized embeddings?*
 - Learn multiple factor embeddings, each of which preserves its desired factor

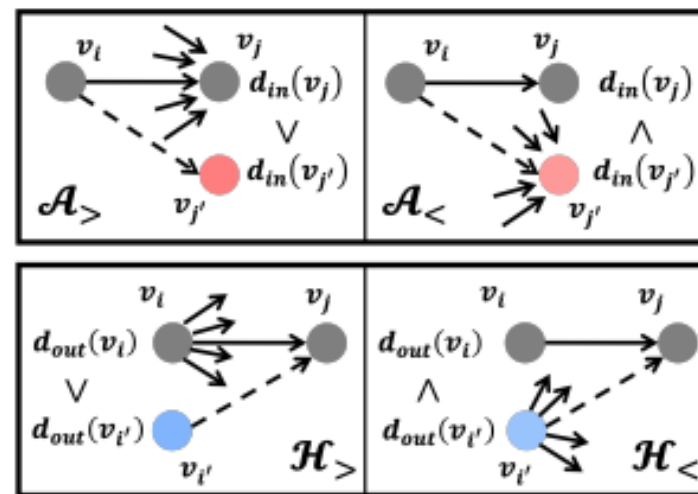
Overview of ODIN




(STAGE 1) Factor Modeling



(STAGE 2) Negative Sampling



(STAGE 3) Disentangled Embedding Learning




(STAGE 1) Factor Modeling

□ Model the formation of each directed edge (v_i, v_j) based on **six node factors** grouped as follows:

1. Authority factors

- (a) The target v_j 's authority status (**a-target**) and (b) the source v_i 's bias toward authorities (**a-source**)
- They together model a bias related to target v_j 's authority status (i.e., in-degree)




(STAGE 1) Factor Modeling (cont'd)

□ Model the formation of each directed edge (v_i, v_j) based on **six node factors** grouped as follows:

2. Hub factors

- (a) The source v_i 's hub status (**h-source**) and (b) the target v_j 's bias toward hubs (**h-target**)
- They together model a bias related to source v_i 's hub status (*i.e.*, out-degree)




(STAGE 1) Factor Modeling (cont'd)

□ Model the formation of each directed edge (v_i, v_j) based on **six node factors** grouped as follows:

3. Interest factors

- (a) The source v_i 's intrinsic property as a source (**i-source**) and (b) the target v_j 's intrinsic property as a target (**i-target**)
- They together model the **pure interest** in forming an edge from v_i to v_j **after removing degree-related biases**



(STAGE 1) Factor Modeling (cont'd)


- Represent a node v_i as **six factor sub-embeddings**
- Concatenate the three factor sub-embeddings as a source/target

■ Source role: $a_i^{src}, h_i^{src}, i_i^{src}$

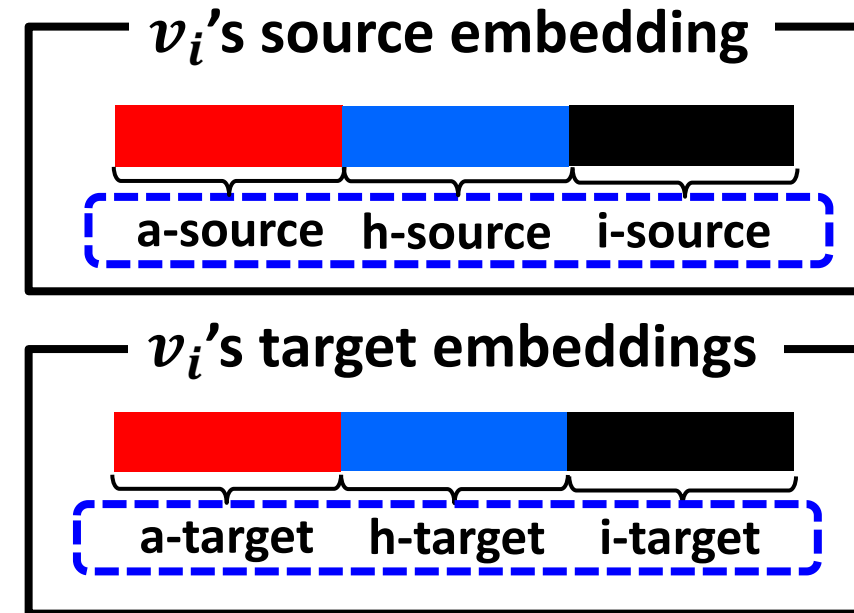
■ Target role: $a_i^{tar}, h_i^{tar}, i_i^{tar}$

■ $s_i = a_i^{src} \oplus h_i^{src} \oplus i_i^{src}$,

■ $t_i = a_i^{tar} \oplus h_i^{tar} \oplus i_i^{tar}$



Six Node Factors



Disentangled Source/Target Embeddings

(STAGE 1) Factor Modeling (cont'd)

□ Compute **three factor scores** based on the six factor sub-embeddings

- Represent how much (a) the authority factor, (b) the hub factor, and (c) the interest factor affect the formation of the directed edge (v_i, v_j)

$$(a) s_{ij}^{auth} = \mathbf{a}_i^{src} \circ \mathbf{a}_j^{tar} \quad (b) s_{ij}^{hub} = \mathbf{h}_i^{src} \circ \mathbf{h}_j^{tar} \quad (c) s_{ij}^{int} = \mathbf{i}_i^{src} \circ \mathbf{i}_j^{tar}$$


□ Compute the **overall edge score by adding the three factor scores**

- Represent the likelihood of the formation of a directed edge (v_i, v_j)

$$s_{ij}^{edge} = s_{ij}^{auth} + s_{ij}^{hub} + s_{ij}^{int} (= \mathbf{s}_i \circ \mathbf{t}_i)$$

(STAGE 2) Negative Sampling

- For each existent edge (v_i, v_j) , sample **different types of training instances (i.e., triplets)** for embedding learning




- Add $(v_i, v_j, v_{j'})$ to the sets $A_{>}$ (for type 1) or $A_{<}$ (for type 2)
 - Aid in capturing the influence of the **bias related to the target's in-degree**
 - $A = A_{>} \cup A_{<}$


(STAGE 2) Negative Sampling (cont'd)

- For each existent edge (v_i, v_j) , sample **different types of training instances (i.e., triplets)** for embedding learning

■ Type 3



■ Type 4



Positive
source

Target

Negative
source

- Add $(v_i, v_j, v_{i'})$ to the sets $H_>$ (for type 3) or $H_<$ (for type 4)

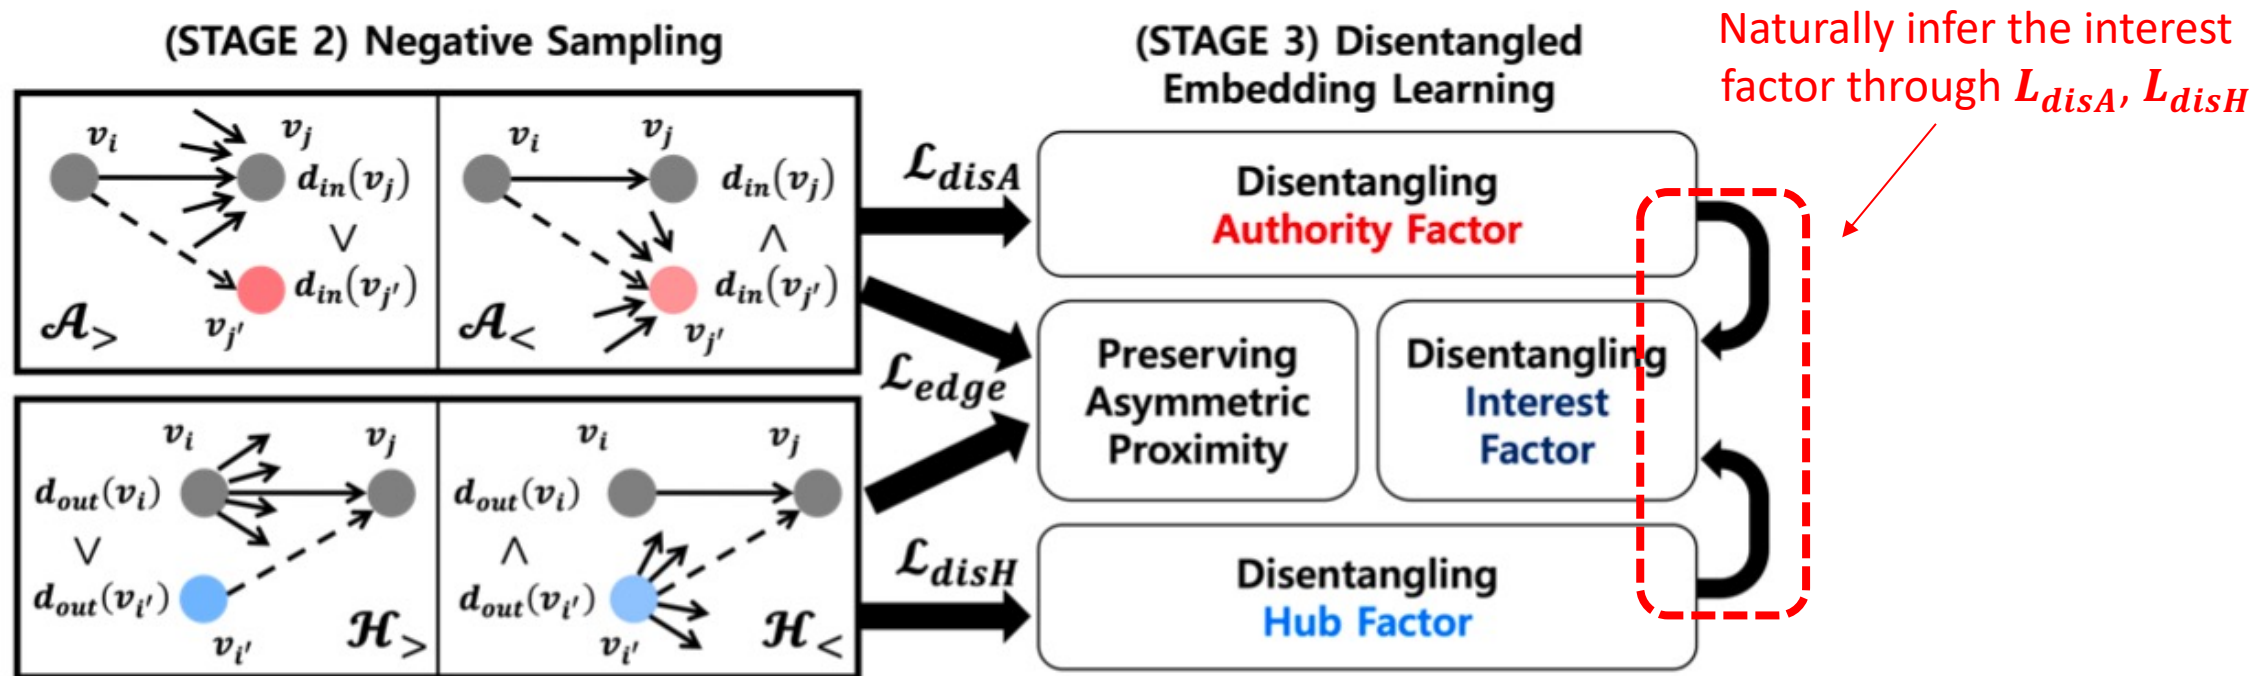
■ Aid in capturing the influence of the **bias related to the source's out-degree**

■ $H = H_> \cup H_<$

(STAGE 3) Disentangled Embedding Learning

□ Learn the disentangled **source and target embeddings** of each node based on the sampled instances via the three objectives

1. L_{edge} : **preserve asymmetric proximity** between nodes in the input network
2. L_{disA} : **disentangle the authority factor** from the other two factors
3. L_{disH} : **disentangle the hub factor** from the other two factors



Loss Function: Multi-Objective Learning

$$L = \boxed{L_{edge}(A \cup H)} + \alpha(\boxed{L_{disA}(A)} + \boxed{L_{disH}(H)})$$

① Preserving
asymmetric proximities

② Disentangling the authority
factor from the others

③ Disentangling the authority
factor from the others

$$\begin{aligned} \textcircled{1} L_{edge}(A \cup H) = & \sum_{(v_i, v_j, v_{j'}) \in A} \text{BPR}(s_{ij}^{edge}, s_{ij'}^{edge}) + \sum_{(v_i, v_j, v_{i'}) \in H} \text{BPR}(s_{ij}^{edge}, s_{i'j}^{edge}) \\ & * \text{BPR}(s_{ij}^{edge}, s_{ij'}^{edge}) = -\log(\sigma(s_{ij}^{edge} - s_{ij'}^{edge})) \end{aligned}$$

→ For $(v_i, v_j, v_{j'})$ in A ,

$s_{ij}^{edge} > s_{ij'}^{edge}$

→ For $(v_i, v_j, v_{i'})$ in H ,

$s_{ij}^{edge} > s_{i'j}^{edge}$

However, L_{edge} alone does not contribute to preserving the desired factor in each factor sub-embedding

Loss Function: Multi-Objective Learning (cont'd)

$$L = \boxed{L_{edge}(A \cup B)} + \alpha \boxed{L_{disA}(A)} + \boxed{L_{disH}(H)}$$

① Preserving
asymmetric proximities

② **Disentangling the authority
factor from the others**

③ Disentangling the authority
factor from the others

$$\textcircled{2} \quad \underline{L_{disA}(A)} = L_{auth}(A_{>}) + L_{auth}(A_{<}) + L_{hub+int}(A_{<})$$

→ For $(v_i, v_j, v_{j'})$ in $A_{>}$, $s_{ij}^{auth} > s_{ij'}^{auth}$

□ The authority status (i.e., in-degree) of v_j is **higher** than that of $v_{j'}$

□ Thus, if authority-factor scores capture the biases towards authorities, as desired,

$s_{ij}^{auth} > s_{ij'}^{auth}$ **should hold!**

Loss Function: Multi-Objective Learning (cont'd)

$$L = \boxed{L_{edge}(A \cup B)} + \alpha \boxed{L_{disA}(A)} + \boxed{L_{disH}(H)}$$

① Preserving
asymmetric proximities

② **Disentangling the authority
factor from the others**

③ Disentangling the authority
factor from the others

$$\textcircled{2} \quad L_{disA}(A) = L_{auth}(A_{>}) + \underline{L_{auth}(A_{<})} + L_{hub+int}(A_{<})$$

➔ For $(v_i, v_j, v_{j'})$ in $A_{<}$, $s_{ij}^{auth} < s_{ij'}^{auth}$

□ The authority status (i.e., in-degree) of v_j is **lower** than that of $v_{j'}$

□ Thus, if authority-factor scores capture the biases towards authorities, as desired,

$s_{ij}^{auth} < s_{ij'}^{auth}$ **should hold!**

Loss Function: Multi-Objective Learning (cont'd)

$$L = \boxed{L_{edge}(A \cup B)} + \alpha(\boxed{L_{disA}(A)} + \boxed{L_{disH}(H)})$$

① Preserving
asymmetric proximities

② **Disentangling the authority
factor from the others**

③ Disentangling the authority
factor from the others

$$\textcircled{2} \ L_{disA}(A) = L_{auth}(A_{>}) + L_{auth}(A_{<}) + \underline{L_{hub+int}(A_{<})}$$

→ For $(v_i, v_j, v_{j'})$ in $A_{<}$,

$$s_{ij}^{hub} + s_{ij}^{int} > s_{ij'}^{hub} + s_{ij'}^{int}$$

□ Even though $s_{ij}^{auth} < s_{ij'}^{auth}$ holds, s_{ij}^{edge} should become **higher** than $s_{ij'}^{edge}$

□ Thus, we can imply the following inequality **$s_{ij}^{hub} + s_{ij}^{int} > s_{ij'}^{hub} + s_{ij'}^{int}$!**

Loss Function: Multi-Objective Learning (cont'd)

$$L = \boxed{L_{edge}(A \cup B)} + \alpha(\boxed{L_{disA}(A)} + \boxed{L_{disH}(H)})$$

① Preserving
asymmetric proximities

② Disentangling the authority
factor from the others

③ Disentangling the hub
factor from the others

$$\textcircled{3} \quad L_{disH}(H) = \underbrace{L_{hub}(H_{>})}_{\text{Hub}} + \underbrace{L_{hub}(H_{<})}_{\text{Hub}} + \underbrace{L_{auth+int}(H_{<})}_{\text{Auth+Int}}$$

For $(v_i, v_j, v_{i'})$ in $H_{<}$, $s_{ij}^{hub} < s_{i'j}^{hub}$

For $(v_i, v_j, v_{i'})$ in $H_{>}$,

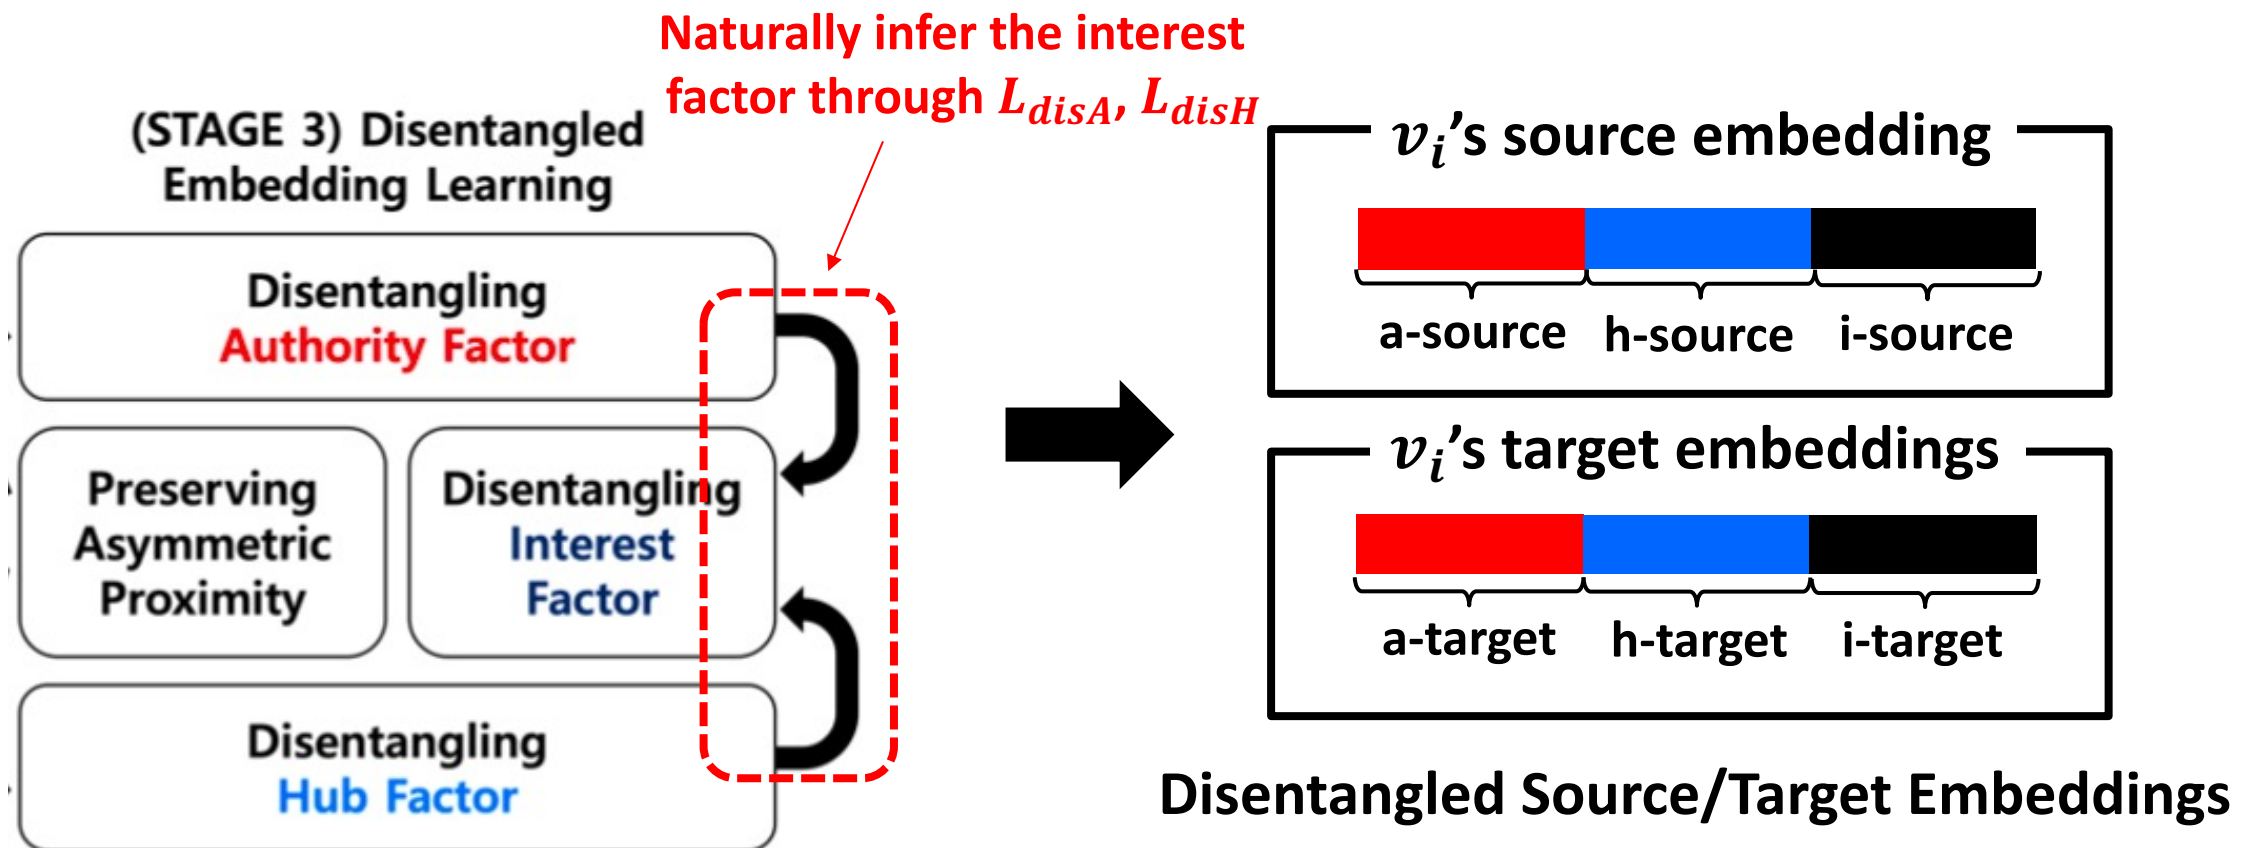
$$s_{ij}^{hub} > s_{i'j}^{hub}$$

For $(v_i, v_j, v_{i'})$ in $H_{<}$,

$$s_{ij}^{auth} + s_{ij}^{int} > s_{i'j}^{auth} + s_{i'j}^{int}$$

Final Embeddings

- Sub-embeddings capture **degree-related biases and interest separately**
- Thus, final embeddings are **robust to the shifts in degree distributions**



Experimental Settings

□ Datasets

Datasets	GNU	Wiki	JUNG	Ciao
Nodes	6,301	7,115	6,120	4,658
Edges	20,777	103,689	50,535	40,133
Reciprocity	0.00%	5.64%	0.90%	34.90%
Types	P2P	Election	Software	Trust

□ Nine competitors

■ 2 undirected NE methods

□ DeepWalk [KDD'14]

□ Node2Vec [KDD'16]

■ 7 directed NE methods

□ APP [AAAI'17]

□ ATP [AAAI'19]

□ NERD [ECML-PKDD'19]

□ GVAE [CIKM'19]

□ DiGCN [NeurIPS'20]

□ MagNet [NeurIPS'21]

□ DGGAN [AAAI'21]

Non-ID Settings

□ Design **two types of non-ID settings** by splitting the edges in an input network into training and test sets **with different degree distributions**

1. Non-ID (in), where **in-degree distributions** are different

□ Each edge (v_i, v_j) is sampled into the test set with $p_{ij}^{in} \propto d_{in}(v_j)^k$

2. Non-ID (out), where **out-degree distributions** are different

□ Each edge (v_i, v_j) is sampled into the test set with $p_{ij}^{out} \propto d_{out}(v_i)^k$


□ Control **the level of distributional shifts by using k**

■ When k is 0, test edges are randomly sampled (i.e., **ID setting**)

■ When k is -1, test edges are sampled inversely proportional to the out-degree of the source nodes or in-degree of the target nodes (i.e., **Shifts are strong**)

Evaluation Task: Link Prediction (LP)

□ How accurately we can predict the directed edges removed from the input directed network?



□ Evaluation protocol

- Consider the **existent edges** as **positive** examples
- Perform two LP tasks, which depend on how we sample the **negative** examples
 - **Uniform LP** (U-LP): consider the **non-existent edges** sampled uniformly at random as **negative** examples
 - **Biased LP** (B-LP): consider the **edges with the opposite directions to (unidirectional) positive examples** as **negative** examples
- Measure **classification accuracy** using *area under curve* (AUC)

Questions to Be Answered

- ❑ RQ1: Does ODIN **outperform its competitors** under distributional shifts in degree distributions?
- ❑ RQ2: How **robust is ODIN under various levels** of distributional shifts in degree distributions?
- ❑ RQ3: Is **factor disentanglement** effective in ODIN?
- ❑ RQ4: How **sensitive is ODIN to its hyperparameters**?

Note: k is fixed to -1 for **RQ1, RQ3, and RQ4** (in $p_{ij}^{in} \propto d_{in}(v_j)^k$ or $p_{ij}^{out} \propto d_{out}(v_i)^k$)

Results for RQ1

Comparison with nine competitors

(a) Non-ID (in)

Datasets	Tasks	Undirected NE		Directed NE							ODIN
		DeepWalk	Node2Vec	APP	GVAE	NERD	ATP	DiGCN	MagNet	DGGAN	
GNU	U-LP	0.593±0.005	0.587±0.004	0.675±0.003	0.675±0.003	0.683±0.008	0.731±0.003	0.729±0.001	<u>0.742±0.001</u>	0.722±0.003	0.760±0.004
	B-LP	0.648±0.006	0.621±0.010	0.700±0.006	0.748±0.013	0.838±0.004	<u>0.910±0.002</u>	0.878±0.003	0.900±0.004	0.901±0.003	0.924±0.001
Wiki	U-LP	0.806±0.001	0.804±0.002	0.795±0.001	0.820±0.005	0.828±0.001	0.827±0.002	0.729±0.002	0.865±0.001	<u>0.890±0.001</u>	0.905±0.001
	B-LP	0.852±0.002	0.855±0.007	0.637±0.008	0.901±0.012	0.915±0.002	0.954±0.001	0.862±0.002	0.928±0.001	<u>0.963±0.001</u>	0.973±0.001
JUNG	U-LP	0.725±0.005	0.777±0.006	0.741±0.002	0.820±0.003	0.784±0.006	0.864±0.001	0.817±0.004	0.816±0.002	<u>0.879±0.003</u>	0.884±0.002
	B-LP	0.810±0.005	0.861±0.005	0.772±0.005	0.902±0.006	0.883±0.005	0.961±0.001	0.926±0.001	0.891±0.003	<u>0.964±0.002</u>	0.969±0.001
Ciao	U-LP	0.776±0.004	0.778±0.002	0.846±0.001	0.841±0.002	0.857±0.002	0.846±0.002	0.641±0.004	0.847±0.001	<u>0.886±0.001</u>	0.892±0.001
	B-LP	0.688±0.005	0.725±0.006	0.768±0.002	0.797±0.004	0.869±0.006	0.887±0.003	0.751±0.006	0.873±0.004	<u>0.912±0.003</u>	0.914±0.003

(b) Non-ID (out)

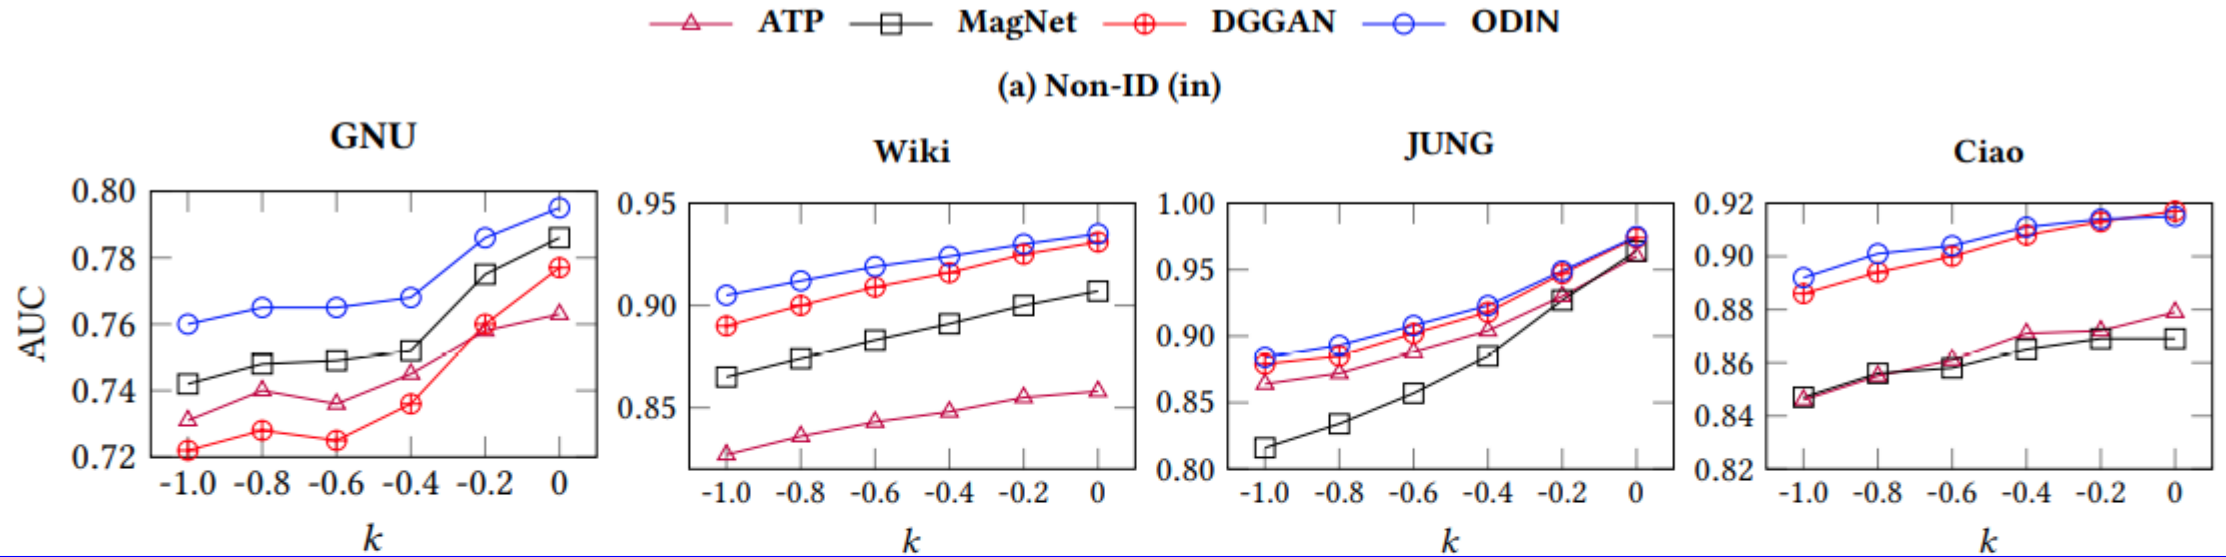
Datasets	Tasks	Undirected NE		Directed NE							ODIN
		DeepWalk	Node2Vec	APP	GVAE	NERD	ATP	DiGCN	MagNet	DGGAN	

- Best competitors change depending on tasks, datasets, and non-ID settings
- ODIN is **effective compared to all the competitors** in addressing the OOD generalization problem against degree-related distributional shifts on directed NE

	B-LP	0.635±0.011	0.695±0.004	0.597±0.004	0.684±0.008	0.750±0.004	0.867±0.003	0.836±0.003	0.827±0.003	<u>0.871±0.002</u>	0.883±0.003
--	------	-------------	-------------	-------------	-------------	-------------	-------------	-------------	-------------	--------------------	--------------------

Results for RQ2

□ Effect of k on the link prediction performance



- In the ID setting (i.e., $k = 0$), the AUCs of ODIN are comparable to or higher than that of the strongest competitors
- ODIN shows the smallest accuracy degradation and, accordingly, the accuracy gain of ODIN against competitors steadily increases
- The results indicate that ODIN obtains OOD generalized embeddings robust to degree-related distributional shifts

Results for RQ3-1

□ RQ3-1: Is each of two disentanglement losses effective in ODIN?

■ ODIN_A vs ODIN_{disA}

Datasets	Tasks	ODIN _A	ODIN _{disA}
GNU	U-LP	0.632±0.005	0.763±0.004
	B-LP	0.704±0.010	0.927±0.001
Wiki	U-LP	0.842±0.002	0.896±0.001
	B-LP	0.918±0.001	0.965±0.001
JUNG	U-LP	0.825±0.004	0.878±0.003
	B-LP	0.929±0.003	0.966±0.002
Ciao	U-LP	0.820±0.003	<u>0.890±0.001</u>
	B-LP	0.788±0.009	<u>0.912±0.002</u>

ODIN_A: only uses the edge loss based on A

ODIN_{disA}: additionally uses the *disA* loss based on A

■ ODIN_H vs ODIN_{disH}

Datasets	Tasks	ODIN _H	ODIN _{disH}
GNU	U-LP	0.604±0.010	0.678±0.001
	B-LP	0.669±0.015	0.820±0.010
Wiki	U-LP	0.793±0.007	<u>0.898±0.001</u>
	B-LP	0.863±0.011	<u>0.968±0.001</u>
JUNG	U-LP	0.714±0.006	0.884±0.002
	B-LP	0.830±0.004	0.970±0.001
Ciao	U-LP	0.853±0.001	0.886±0.001
	B-LP	0.867±0.005	0.909±0.002

ODIN_H: only uses the edge loss based on H

ODIN_{disH}: additionally uses the *disH* loss based on H

■ Each of the disentanglement losses is effective in obtaining embeddings robust to distributional shifts in degree distributions

Results for RQ3-2

□ RQ1-2: Is jointly using the both losses effective in ODIN?

(a)
Non-ID (in)

Datasets	Tasks	ODIN _A	ODIN _{disA}	ODIN _H	ODIN _{disH}	ODIN
GNU	U-LP	0.632±0.005	0.763±0.004	0.604±0.010	0.678±0.001	<u>0.760±0.004</u>
	B-LP	0.704±0.010	0.927±0.001	0.669±0.015	0.820±0.010	<u>0.924±0.001</u>
Wiki	U-LP	0.842±0.002	0.896±0.001	0.793±0.007	<u>0.898±0.001</u>	0.905±0.001
	B-LP	0.918±0.001	0.965±0.001	0.863±0.011	<u>0.968±0.001</u>	0.973±0.001
JUNG	U-LP	0.825±0.004	0.878±0.003	0.714±0.006	0.884±0.002	0.884±0.002
	B-LP	0.929±0.003	0.966±0.002	0.830±0.004	0.970±0.001	<u>0.969±0.001</u>
Ciao	U-LP	0.820±0.003	<u>0.890±0.001</u>	0.853±0.001	0.886±0.001	0.892±0.001
	B-LP	0.788±0.009	<u>0.912±0.002</u>	0.867±0.005	0.909±0.002	0.914±0.003

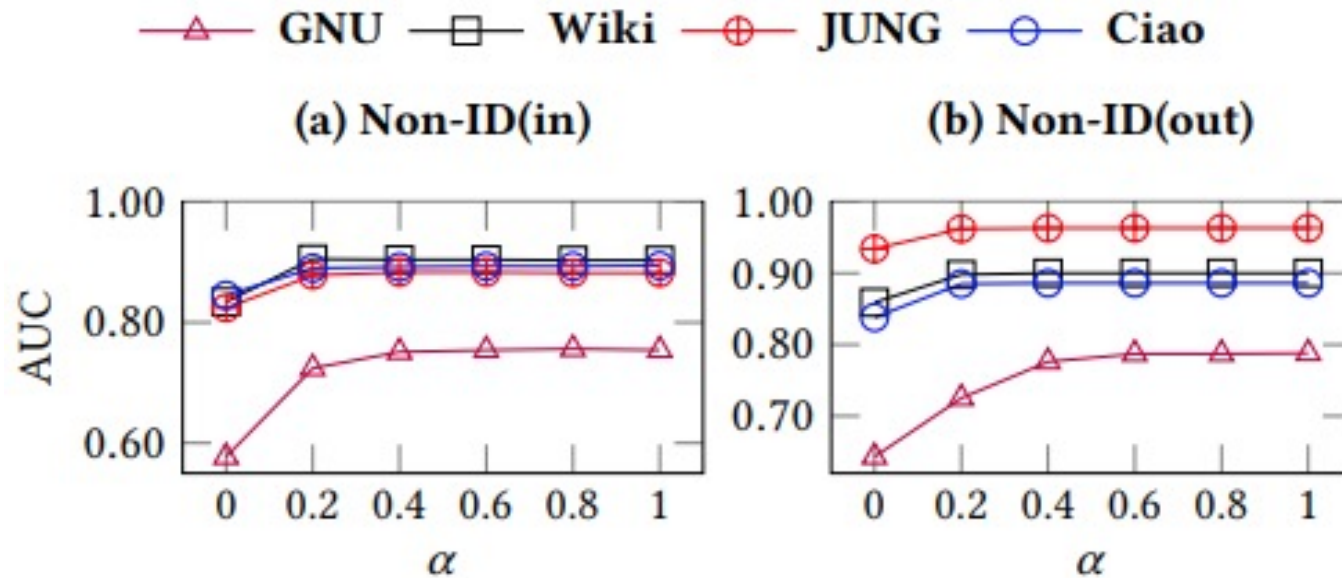
■ Superiority between ODIN_{disA} and ODIN_{disH} varies depending on datasets

■ ODIN outperforms ODIN_{disA} and ODIN_{disH} in most cases

□ That is, ODIN can selectively adopt the factor(s) beneficial in each dataset, thereby improving the robustness of embeddings in all datasets

Results for RQ4

□ How the parameter α affects the accuracy of ODIN



$$L = L_{edge}(A \cup B) + \alpha(L_{disA}(A) + L_{disH}(H))$$

- AUCs of ODIN steadily increase until α reaches 0.4 and then the AUCs converge
- ODIN is not highly sensitive to the weight for factor disentanglement

Conclusions

- We pointed out that the existing directed NE methods **face difficulties in effectively addressing the OOD generalization problem**
- We proposed **ODIN**, which **models multiple factors in the formation of directed edges and learns nodes' multiple factor sub-embeddings**
- Through extensive experiments, we showed clearly the effectiveness of our strategies for **factor modeling** and **disentangled embedding learning**